

Overview

This assignment focuses on a few fundamental concepts in signal processing. In the previous assignment, the signals were discrete events like photons. Here, we will look at continuous signals that have structure that extends over time with additive noise.

Learning objectives

- explain and illustrate the discrete representation of a continuous signal
- explain aliasing and the Nyquist frequency
- explain the delta and step functions
- write functions to synthesize signals
- write functions that accept functions as arguments
- estimate the energy and power of a signal
- compute the signal to noise ratio
- generate signals with different levels of additive noise
- design closed analysis/synthesis loops for testing correctness

Readings

- The following material provides both background and additional context. It is linked in the Canvas page for this assignment. Refer to these for a more detailed explanation and formal presentation of the concepts in the exercises. Note that the readings contain a lot more material than what's needed for the exercises. Feel free to skim or skip the parts that aren't directly relevant.
- Dostert, D. B. (1992). Sensory Ecology. Chapter 5-3 Signal Processing.
- Proakis and Salehi (2008) Signal Processing for Communications: Chapters 1 and 2.

Exercises

```
In [ ]: import numpy as np
import math

from matplotlib import pyplot as plt

In [ ]: import functions as f # import my functions from A3b and A3a (functions in A3a are shown below)
```

1. Continuous signals and sampling

Discrete sampling of continuous signals has some implications that are important to appreciate. The first is that the sampled signal is only a representation of the underlying continuous signal and it doesn't necessarily capture all the information. It is easy to visualize this by making a plot that overlays the discrete samples on the continuous function.

1a. Sampled functions

Write a function

```
def plot_sampled_function(gsin: Callable[[float], float], fs: float = 1, tlim: tuple = (0, 2*np.pi), tscale: float = 1, tunits: str = "secs"):
```

which plots a function `g` with overlaid samples at sampling frequency `fs` over range `tlim`, which is scaled by `tscale` and is plotted in units of `units`. The function should be plotted as a line, and the samples should be overlaid as a stem plot.

Looking ahead to the next module on sound localization, we want to start getting used to thinking in terms of the time scale of the waveform and use examples that are more relevant to biological perception. Instead of the defaults (which are sensible for a generic function), use `tscale=1/900` and `units="msecs"` for your examples. Note that this time scale only applies to the plot – time as a function argument should be in seconds. This is to have a clean separation between the information and the display of the information.

Get two examples using a sine and gammatone functions. You can re-use your code from A3b.

```
In [ ]: def plot_sampled_function(gnp: Callable[[float], float], fs: float = 0.01, fs_float: float = 0, tlim: tuple = (0, 2*np.pi), tscale: float = 1, tunits: str = "secs"):
    # function over tlim
    T = np.arange(tlim[0], tlim[1], tscale, f)
    F = g(F)

    # samples
    S = np.arange(tlim[0], tlim[1], tscale, f)
    S = f(S)

    # plot
    plt.plot(F, F, c='b', label='signal')
    plt.stem(S, S, label='samples')
    plt.title('Sampled Function')
    plt.xlabel('Time (units)')
    plt.ylabel('Amplitude')
    plt.legend()
    plt.show()
```

```
In [ ]: plot_sampled_function(f.sinewave, f=0.001, fs=0.005, tlim=(0, 1), tscale=1/900, tunits="msecs")
```

1b. The Nyquist frequency and aliasing

The **Nyquist frequency** is defined as  $f_s/2$ , i.e. one half the sampling frequency. It is the upper bound on the signal frequency  $f_s$ .

Aliasing occurs when the frequency of the signal is greater than the sampling frequency, i.e.  $f > f_s$ . In this case, sampling of a periodic function like sine can result in the appearance of sampling a function of much lower frequency than what's actually there. To avoid aliasing artifacts when you digitally sample signals, you need to filter out all frequencies higher than the Nyquist frequency.

Use your function above to illustrate different types of sampling phenomena for these conditions:

1. A sine wave below Nyquist at a frequency that shows a few samples per period which unevenly distributed.
2. sine at Nyquist
3. cosine at Nyquist
4. cosine sampled above Nyquist frequency that clearly shows aliasing

Note that there are other common usages of the term "aliasing". One is a technical sense that we have just explained. Another is more colloquial and is used to describe any situation where the samples don't reflect the true underlying pattern. It is commonly used to describe plotting a high frequency waveform where the discrete samples, even though there is no aliasing in a technical sense, show a more jagged structure than the actual analog pressure waveform, which is smooth. Another common usage is when the details of a plot don't align well with the pixels.

```
In [ ]: # sine wave below Nyquist frequency
plot_sampled_function(gnp.sine, f=0.001, fs=0.005, tlim=(0, 2*np.pi))
```

```
In [ ]: # cosine at Nyquist frequency
plot_sampled_function(gnp.cos, f=0.1, fs=0.2, tlim=(0, 2*np.pi))
```

```
In [ ]: # cosine above Nyquist frequency
plot_sampled_function(gnp.cos, f=1, fs=0.3, tlim=(0, 2*np.pi))
```

2. Signals

We have used functions like sine wave, Gaussian, Gabor, and gammatone. Here we add two more functions to our library that will be useful for generating signals.

The delta function

The Dirac delta function is used to to model an impulse or discrete event as a brief impulse of energy. The delta function is zero everywhere except at  $t = 0$

$$\delta(t) = \begin{cases} \text{undefined} & t = 0 \\ 0 & t \neq 0 \end{cases}$$

The value at zero is unbounded and undefined, but the integral is one

$$\int_{-\infty}^{\infty} \delta(t) dt = 1$$

We also have the property

$$\int_{-\infty}^{\infty} f(t)\delta(t - \tau) dt = f(\tau)$$

Another way to think about this is that  $\delta(t - \tau)$  is zero everywhere except at  $t = \tau$ . At that (infinitesimal) point,  $f(t - \tau)$  is constant and so multiplies the integral, which is one.

The unit step function

The unit step function (also called the Heaviside step function) is used to indicate a constant signal that starts at  $t = 0$ . It is defined by

$$u(t) = \begin{cases} 1 & t \geq 0 \\ 0 & t < 0 \end{cases}$$

2a. Delta and step functions

Write two functions `delta(t; fs)` and `u(t)` to implement the delta and unit step functions. To use these functions to generate signals, which you can then sample, define them so that they accept a continuous time value.

In the case of the delta function, we need define what we mean by  $\tau = 0$ , since that will depend on the sampling frequency. To see why, note that we can model the sampling process as the integration of a function over the sample period so

$$y = \int_{t-\Delta t}^{t+\Delta t} f(t) dt$$

where  $y$  is the sample value,  $\Delta t = 1/f_s$  is the sample period, and  $t$  is the sample time which is centered on the period. This means that if an impulse falls anywhere within a given sample period, the value of that sample will be one. This is an idealized model. Real-world impulses are rarely within the bounds of a sample.

```
In [ ]: def delta(t=0, fs: float = 1):
    # function of tlim
    c = 0
    T = np.arange(t)
    F = np.arange(t)
    if T.size > 1:
        return np.array([1 if abs(t-c) <= abs(c-1/(fs*2)) else 0 for t in T])
    else:
        return 1 if abs(t) <= abs(c-1/(fs*2)) else 0
```

```
In [ ]: T = np.arange(-1, 5, 0.01)
F = delta(T, fs=100)
plt.plot(T, F)
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.title('Dirac Delta Function')
plt.show()
```

```
In [ ]: def u(t=0):
    # function of tlim
    c = 0
    T = np.arange(t)
    F = np.array([1 if t >= c else 0 for t in T])
    return np.array([1 if t >= c else 0 for t in T]) if T.size > 1 else 1 if t >= c else 0
```

```
In [ ]: T = np.arange(-1, 5, 0.01)
F = u(T)
plt.plot(T, F)
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.title('Unit Step Function')
plt.show()
```

2b. gensignal

Write a function `gensignal(t; g, T, T1, T2)` that generates the values at time `T` of a signal defined by function `g`, which should be function of time. Other arguments to `g` can be specified upon definition, e.g.

$$x = \text{gensignal}(t, t \rightarrow \text{gammatone}(t; f=100), t=0.025, T=0.1)$$

The signal should be delayed by `T` and have duration limited to `T1`, i.e. it has value `g(t - T)` and is zero for  $t < T$  and  $t > T + T1$ . Note that  $T + T1$  is an exclusive limit, because the sample times are centered on the sample periods. For example, a unit step function for  $f_s = 1$ ,  $T = 0$ , and  $T1 = 2$  will have values one only at times 0 and 1 with sample periods that extend from  $-0.5/f_s$  to  $1.5/f_s$ .

```
In [ ]: def gensignal(t, g: Callable[[float], float], tau: float = 0, T: float = 0):
    # generate signal
    T = np.arange(t)
    F = g(T-Tau)
    return np.array([0 if T[i] <= tau or T[i] >= T+tau else F[i] for i in range(len(T))]) if T.size > 1 else 0 if T <= tau or T >= T+tau else F
```

```
In [ ]: T = np.arange(-1, 10, 0.01)
F = gensignal(T, g=lambda t: f.sinewave(t), tau=2, T=2)
plt.plot(T, F)
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.title('gensignal Function')
plt.show()
```

3. Noise and SNR

It is useful to have a generic way of describing the detectability of a signal. In the Signal Detection assignment, we characterized this with the probability distributions of the events and the noise. There, the events were discrete and we assumed they occurred within a sample. For signals that extend over time, it is common to use the signal to noise ratio.

The signal to noise ratio ("SNR") is simply the power of the signal divided by the power of the noise. So, naturally you ask, "What is power?" That's a deep and complex question, but in the case of signals, power is average energy over a period. The energy of a signal  $x(t)$  is defined as

$$E_x = \sum_{n=1}^N |x[n]|^2$$

Here we will assume that  $n$  runs over the extent of the signal, e.g. a sound of length  $N$ . Thus, the energy is the same as the squared norm  $\|x\|^2$  we used in A1b. The power of  $x$  is then

$$P_x = \frac{1}{N} \sum_{n=1}^N |x[n]|^2 = \sigma_x^2$$

where  $[n]$  indicates 1-based array indexing rather than a discrete time value of the function  $x(t)$ . Note that for zero mean, the power is equivalent to variance.

For a signal in additive noise

$$y(t) = x(t) + n(t)$$

the SNR is simply

$$\frac{P_x}{P_n}$$

It is almost always expressed on a logarithmic scale in units of decibels (dB)

$$\text{dB SNR} = 10 \log_{10}(P_x/P_n) = 20 \log_{10}(\sigma_x/\sigma_n)$$

Note that this implicitly assumes that we know the extent of the signal (to calculate  $P_x$ ) or that it is **stationary** in time, i.e. the signal's structure doesn't change over time and extends throughout the period of analysis. Structure could be described by the frequency content or by a probability distribution.

Peak Signal Noise Ratio

A related concept, which we won't use here, but is used more often in image processing, is peak signal to noise ratio or PSNR. Many signals have limited extent which we don't know a priori, e.g. a feature in an image. In this case, it makes sense to use the maximum value to approximate the best (or peak) SNR, i.e. the point where the signal is strongest.

Since we don't know the signal, we also don't know the noise, so a second approximation is to assume that the signals are sparse (rarely occurring). In this case, the power (or variance) of the noise can be approximated with the variance of the observed waveform  $y$ , because we assume it is dominated by the noise. An example in images would be sparse features on a smooth background where the variance would be dominated by the "smooth" background, and so would approximate the underlying noise.

Putting these concepts together (and again assuming zero mean) gives

$$\text{PSNR} = 10 \log_{10} \left( \frac{\max(y(t))^2}{\sigma_y^2} \right) = 20 \log_{10} \left( \frac{\max(y(t))}{\sigma_y} \right)$$

3a. energy, power, and snr

Write functions `energy(x)`, `power(x)`, `snr(Ps, Pn)` for the definitions above.

```
In [ ]: def energy(x):
    # sum all elements to 2nd power and sum all elements
    X = np.array(x)
    return np.sum(np.power(X, 2))
```

```
In [ ]: def power(x):
    X = np.array(x)
    return energy(x)/X.size
```

```
In [ ]: def snr(Ps, Pn):
    return 10 * np.log10(Ps / Pn)
```

3b. Additive signals

Write a function `y = noisysignal(t; g, T, T1, T2, sigma)` that generates a sample at time `T` of a signal plus additive Gaussian noise. Like above, the signal is delayed by `T`, `sigma` specifies the standard deviation of the noise. Show examples with a sine wave, step, and gammatone

```
In [ ]: def noisysignal(t, g: Callable[[float], float], tau: float = 0.25, T: float = 1, sigma: float = 0.05):
    # gensignal for range t
    T = np.arange(t)
    F = g(T-Tau)
    # generate noise
    N = np.random.normal(0, sigma, T.size)
    return np.add(F, N)
```

```
In [ ]: T = np.arange(-1, 5, 0.01)
F = noisysignal(T, g=lambda t: f.sinewave(t), tau=0.25, T=4, sigma=0.05)
plt.plot(T, F)
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.title('noisysignal Function')
plt.show()
```

3c. Noise level specified by SNR

In 3b, we added noise that had a fixed variance. Here we want to generate a signal that has a noise level specified by an SNR. Since the SNR is the average signal power, it depends on the signal. Thus, to calculate the noise level needed to achieve a specified SNR, we need to define a function that accepts an array as input and also the location of the signal in the array.

Write a function

```
def snr(x: np.ndarray, x_range: slice, snr: float) -> np.ndarray
```

which, given array `x`, returns the standard deviation of additive Gaussian noise such adding noise at that level to `x` has an SNR of `snr` dB. The optional argument `x_range` specifies location of the signal, i.e. the range over which to compute the signal power. It should default to the whole signal.

Note that calculating the signal power over the whole waveform when the signal is only present in part of the waveform would lead to a biased result. Why is this? Illustrate this by contrasting the resulting waveforms produced with and without knowledge of signal location.

```
In [ ]: # gen signal
x = np.arange(0, 10*np.pi, 0.01)
F = gensignal(x, g=lambda t: np.sin(t), tau=5, T=8)
# calc snr over entire range and only signal range
print(f'Signal over entire range: {snr2sigm(F, snr=20)}')
print(f'Signal over signal range: {snr2sigm(F, x_range=slice(8, 20), snr=20)}')
```

Signal over entire range: 0.0575450877525985  
Signal over signal range: 0.02899732207783819

3d. Estimating SNR

One of the challenges in developing algorithms for perceptual computations is that we rarely know the ground truth, and we often don't have control over the signal structure or real world conditions. It is therefore useful to develop methods for synthesis, in that spirit, we will "complete the loop" and estimate the SNR by a waveform.

Write a function `extent(y; B=8)` that returns a range from the first to last index where the absolute value of array `y` known threshold `B`, which is specified as a fraction of the maximum absolute value.

Show that it produces the correct index range for a known case, and use it to estimate the SNR for a synthesized signal with known SNR.

```
In [ ]: def extent(y, theta=0.01):
    # estimate threshold
    max = np.max(y)
    Y = np.abs(y)
    if Y >= max:
        t = theta * max
    # return start and end indices
    i = np.where(Y >= t)
    return (i[0], i[-1])
```

```
In [ ]: # find snr for snr of 10
T = np.arange(-1, 10*np.pi, 0.01)
F = gensignal(T, g=lambda t: np.sin(t), tau=5, T=8)
print(f'Estimated snr for snr of 10: {snr2sigm(F, snr=20)}')
```

Estimated snr for snr of 10: 0.05673367535697077

```
In [ ]: # estimate SNR for known signal
T = np.arange(-1, 10*np.pi, 0.01)
F = noisysignal(T, g=lambda t: np.sin(t), tau=5, T=20, sigma=0.0507)
```

# find signal  
s, e = extent(F, theta=0.25)  
print(f'Signal found between indices {s}, {e}')  
# separate signal from noise  
F = F[s:e]  
F = np.concatenate((F[s], F[e-1]))

# calc snr  
snr = snr2sigm(F[s:e], Pn=power(F[s:e]))  
print(f'Estimated SNR: {snr}')  
Signal found between indices [620, 2599]  
Estimated snr for snr of 10: 0.05673367535697077

```
In [ ]: plt.plot(T[s:e], F[s:e])
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.title('Plot of identified signal')
plt.show()
```

4. Grand synthesis

One measure of the quality of your code design is the ease and flexibility of expressing new ideas. To test this, use your functions to synthesize a waveform composed of random, normalized gammatones plus some level of Gaussian noise.

$$s \sim \text{uniform}(0, T) \\ f_i \sim \text{uniform}(f_{min}, f_{max}) H_s$$

Synthesize a several second waveform and export it to a wav file. What does it sound like? Feel free to experiment with different parameters and distributions.

```
In [ ]: # import packages
from scipy.io import wavfile

# set the parameters
Fs = 44100 # Hz
dur = 10 # seconds
filename = "out.wav"
```

```
In [ ]: # times for plotting
T = np.linspace(0, 3, 3*Fs)

# parameters to create gammatones
delays = [np.random.uniform(0, S) for i in range(num)]
freqs = [np.random.uniform(10, 10000) for i in range(num)]

# generate the signals
S = np.zeros(len(T))
for i in range(1, num+1):
    S = S + g(freqs[i], delays[i])

# normalize
S_norm = S / np.max(S)

# write to wav file
wavfile.write(filename, Fs, S_norm)
```

Tests and self checks

You should write tests for your code and make plots to verify that your implementations are correct. After you submit your draft version, take the self check quiz. This will give you feedback so you can make corrections and revisions before you submit your final version. Here are examples of the types of questions you can expect

- plotting different sampled functions given specified ranges and sampling frequencies
- conceptual questions regarding aliasing
- plotting the result of `gensignal` using specified functions and parameters
- calculating energy, power, and SNR for test waveforms
- plotting noisy signals using specified parameters
- demonstrating the extent function
- estimating SNR from a test waveform given the signal range
- providing an example of a synthesized sum of gammatones.

Submission Instructions

Please refer to the Assignment Submission Instructions on canvas under the Pages tab.